


Continuous pprof collection in Go


Profiling thousands of agents running at the Edge

Ozan Sazak - June 2022



Who am I

- ▶ Ozan Sazak
- ▶ Senior CS student at METU
- ▶ Part-time SWE at Edge Delta



Agenda

- ▶ What is Edge Delta?
- ▶ Why do we need continuous profiling?
- ▶ Brief introduction to pprof
- ▶ Examples of pprof usage
- ▶ How does our continuous profiling work?
- ▶ Demo on continuous profiling

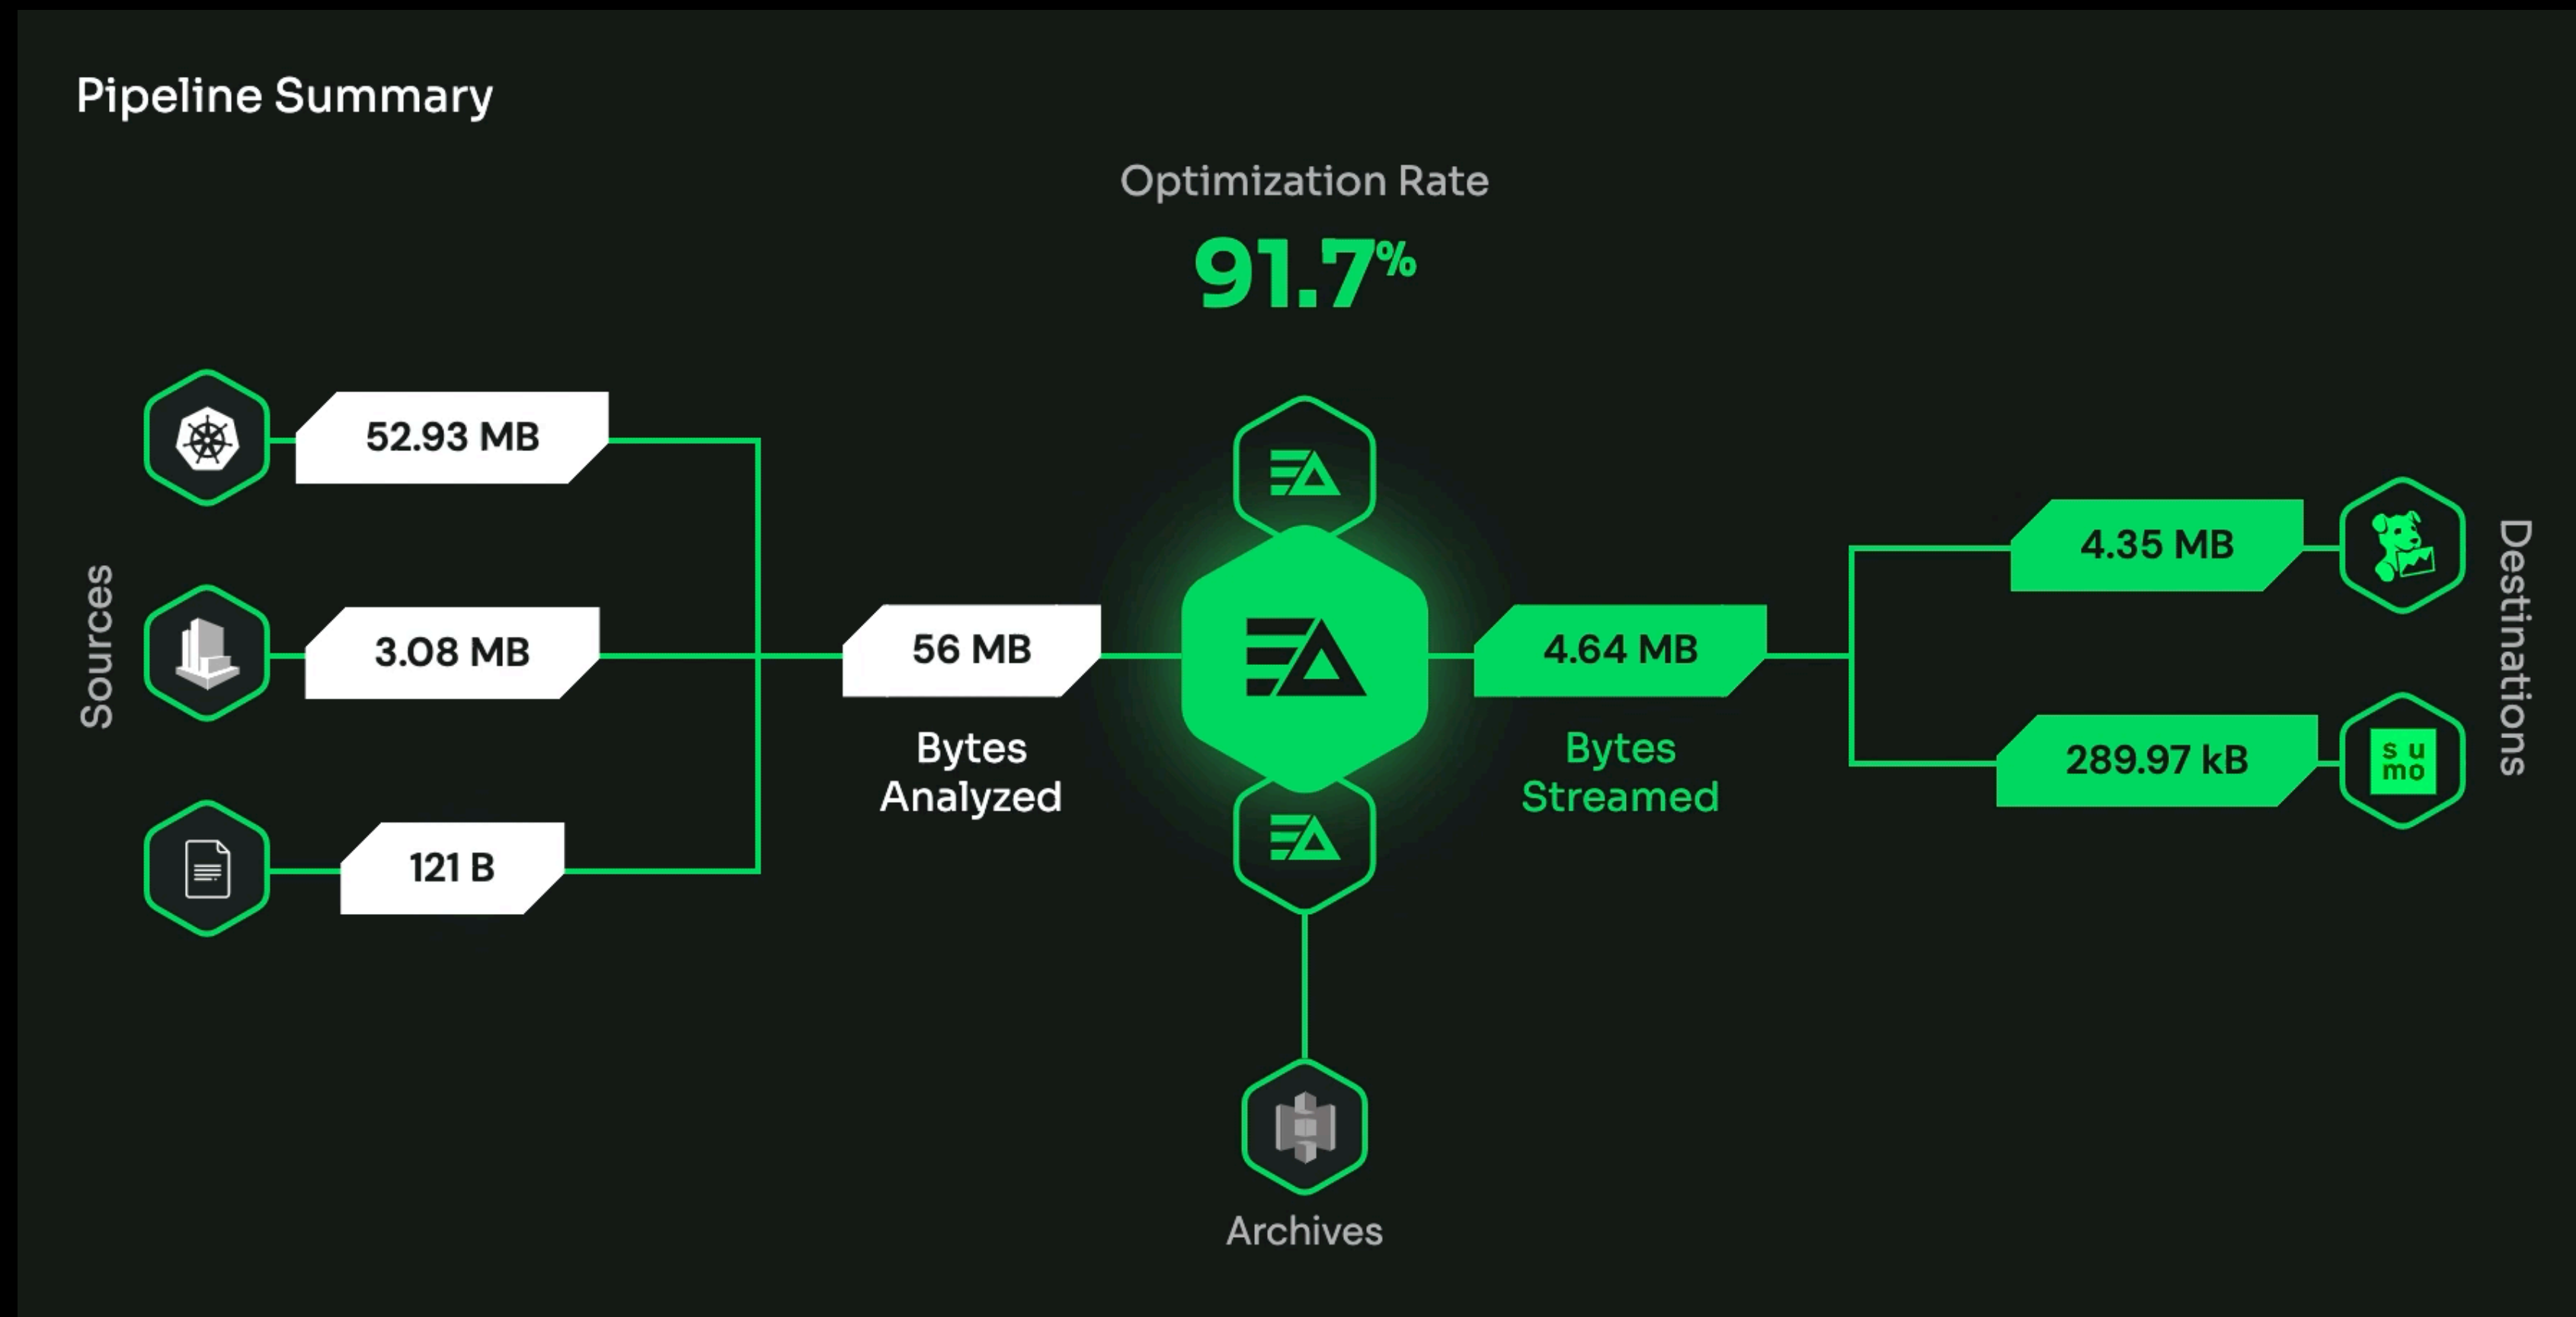
Edge Delta

- ▶ Continuous log observation with thousands of agents running on the edge
- ▶ Agents preprocess logs and filter/tag findings
- ▶ Agents run on customers' machines with a single YAML config
- ▶ (We're actually using it too, to observe our own services :))



Edge Delta

Multiple input sources, processors, and output integrations through agent config



... why do we need continuous profiling?

- ▶ Agents run on customers' **production** machines
 - ▶ They have to be **super performant**
 - ▶ It's **critical** for us to have more visibility on our **bottlenecks**
- ▶ We run nightly performance tests before agent releases
 - ▶ Gives us insight and confidence over agent performance

but...

There are always unforeseen cases on customer environment due to,

- ▶ Different ways of agent configuration
- ▶ Limited resources
- ▶ Millions of files
- ▶ ...

What is profiling?

- ▶ Go is often used in programs with high-performance needs
- ▶ We also face memory leaks and bottlenecks
- ▶ Need to analyze CPU and memory intensive code for optimizations

goroutine	– stack traces of all current goroutines
heap	– a sampling of memory allocations of live objects
allocs	– a sampling of all past memory allocations
threadcreate	– stack traces that led to the creation of new OS threads
block	– stack traces that led to blocking on synchronization primitives
mutex	– stack traces of holders of contended mutexes

The CPU profile is not available as a Profile. It has a special API, the `StartCPUProfile` and `StopCPUProfile` functions, because it streams output to a writer during profiling.

Source: pkg.go.dev/runtime/pprof

Different ways to collect profiles

▶ net/http/pprof

- ▶ Lets us download profiles directly on demand using the default HTTP server
- ▶ /debug/pprof/{goroutine, heap, threadcreate, block, mutex, profile}

▶ runtime/pprof

- ▶ Lets us generate and save profiles inside the code, in runtime

- ▶

```
pprof.Lookup("heap")
```


▶ go test

- ▶ Lets us dump profiles automatically after the test runs

`go test` with `-cpuprofile`

```
func fibonacci(n uint64) uint64 {  
    if n <= 1 {  
        return n  
    }  
    return fibonacci(n-1) + fibonacci(n-2)  
}
```

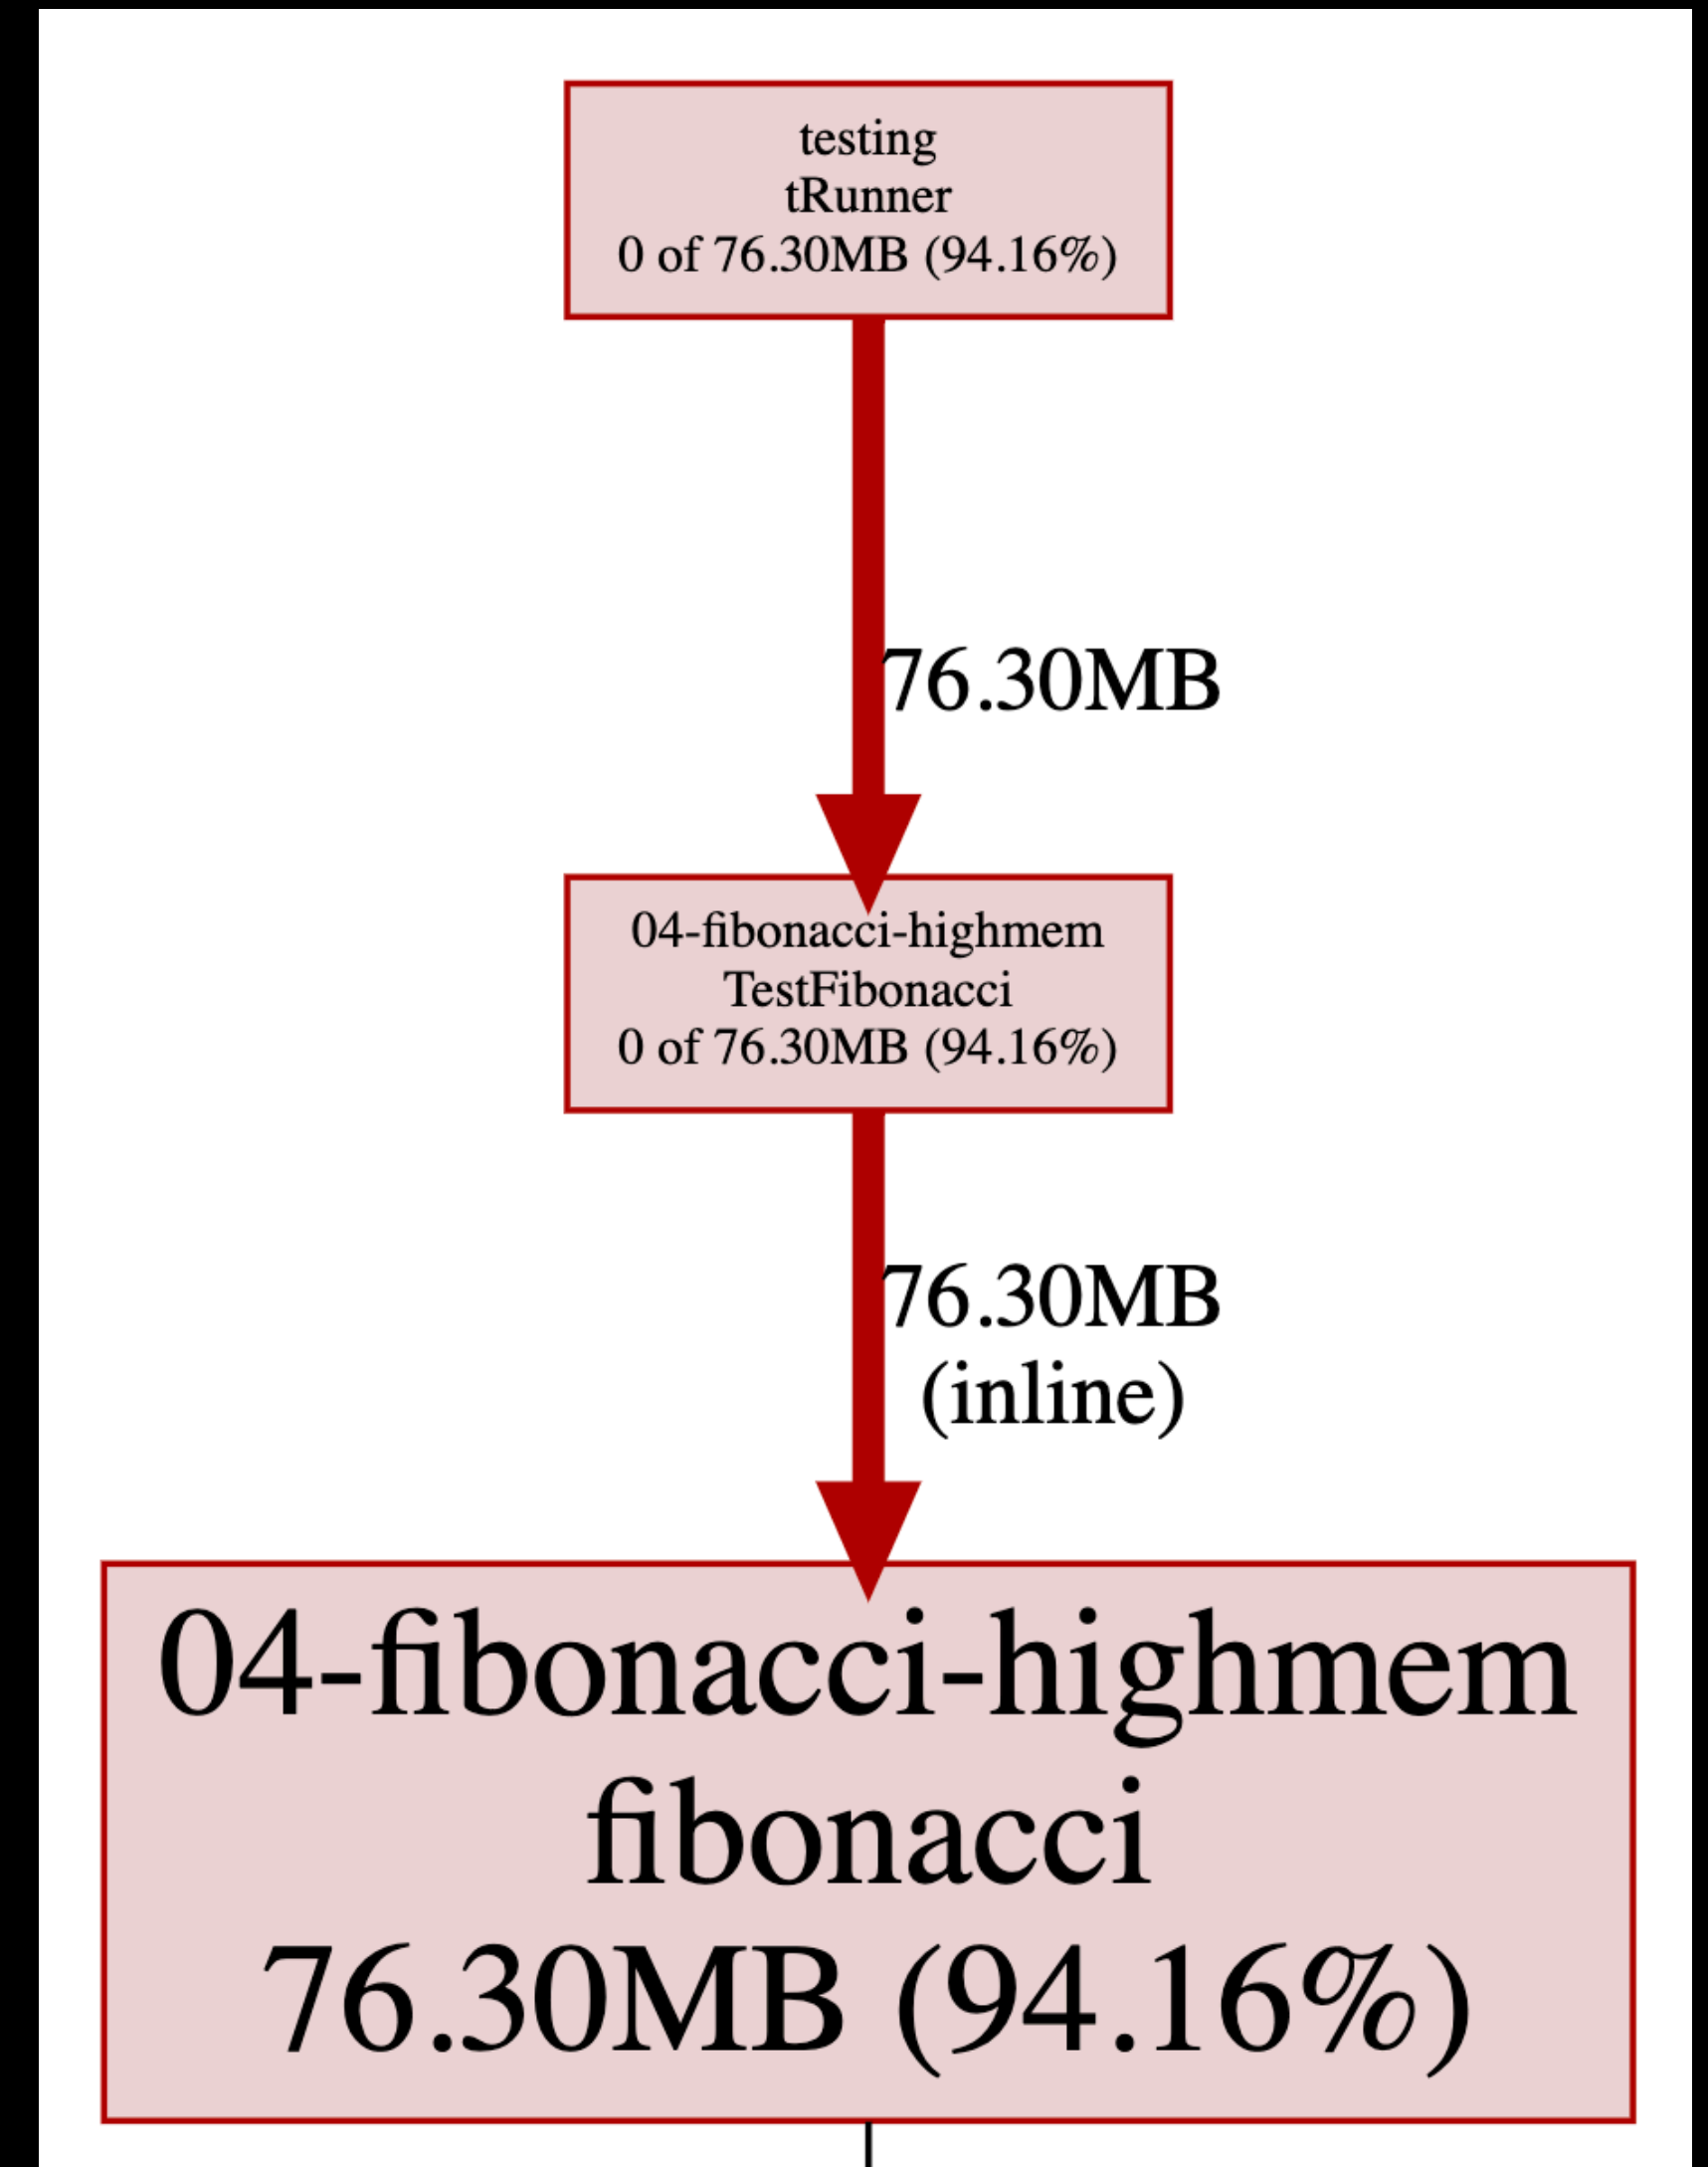
```
> go test -cpuprofile cpu.prof -bench .
```



`go test` with -memprofile

```
func fibonacci(n uint64) uint64 {  
    cache := make([]uint64, n+1)  
    cache[0] = 0  
    cache[1] = 1  
    var i uint64  
    for i = 2; i <= n; i++ {  
        cache[i] = cache[i-1] + cache[i-2]  
    }  
    return cache[n]  
}
```

```
> go test -cpuprofile cpu.prof -memprofile mem.prof -bench .
```



...with pprof HTTP server

```
import _ "net/http/pprof"

func main() {
    go func() {
        log.Println(http.ListenAndServe("localhost:6060", nil))
    }()
    // Do some memory and CPU intensive work here...
}
```

```
> go tool pprof http://localhost:6060/debug/pprof/heap
```

Some useful commands for `go tool pprof`

- ▶ **topN**: Shows the top N samples in the profile
- ▶ **web**: Generate a graph of the profile data in SVG format and opens it on the browser
- ▶ **pdf**: Generates a PDF file with the same graph as the **web** command does
- ▶ **png**: Generates a PNG file with the same graph as the **web** command does
- ▶ **list func**: Shows the source code of the **func** with the flat and cum metrics side by side

How does continuous profiling work?

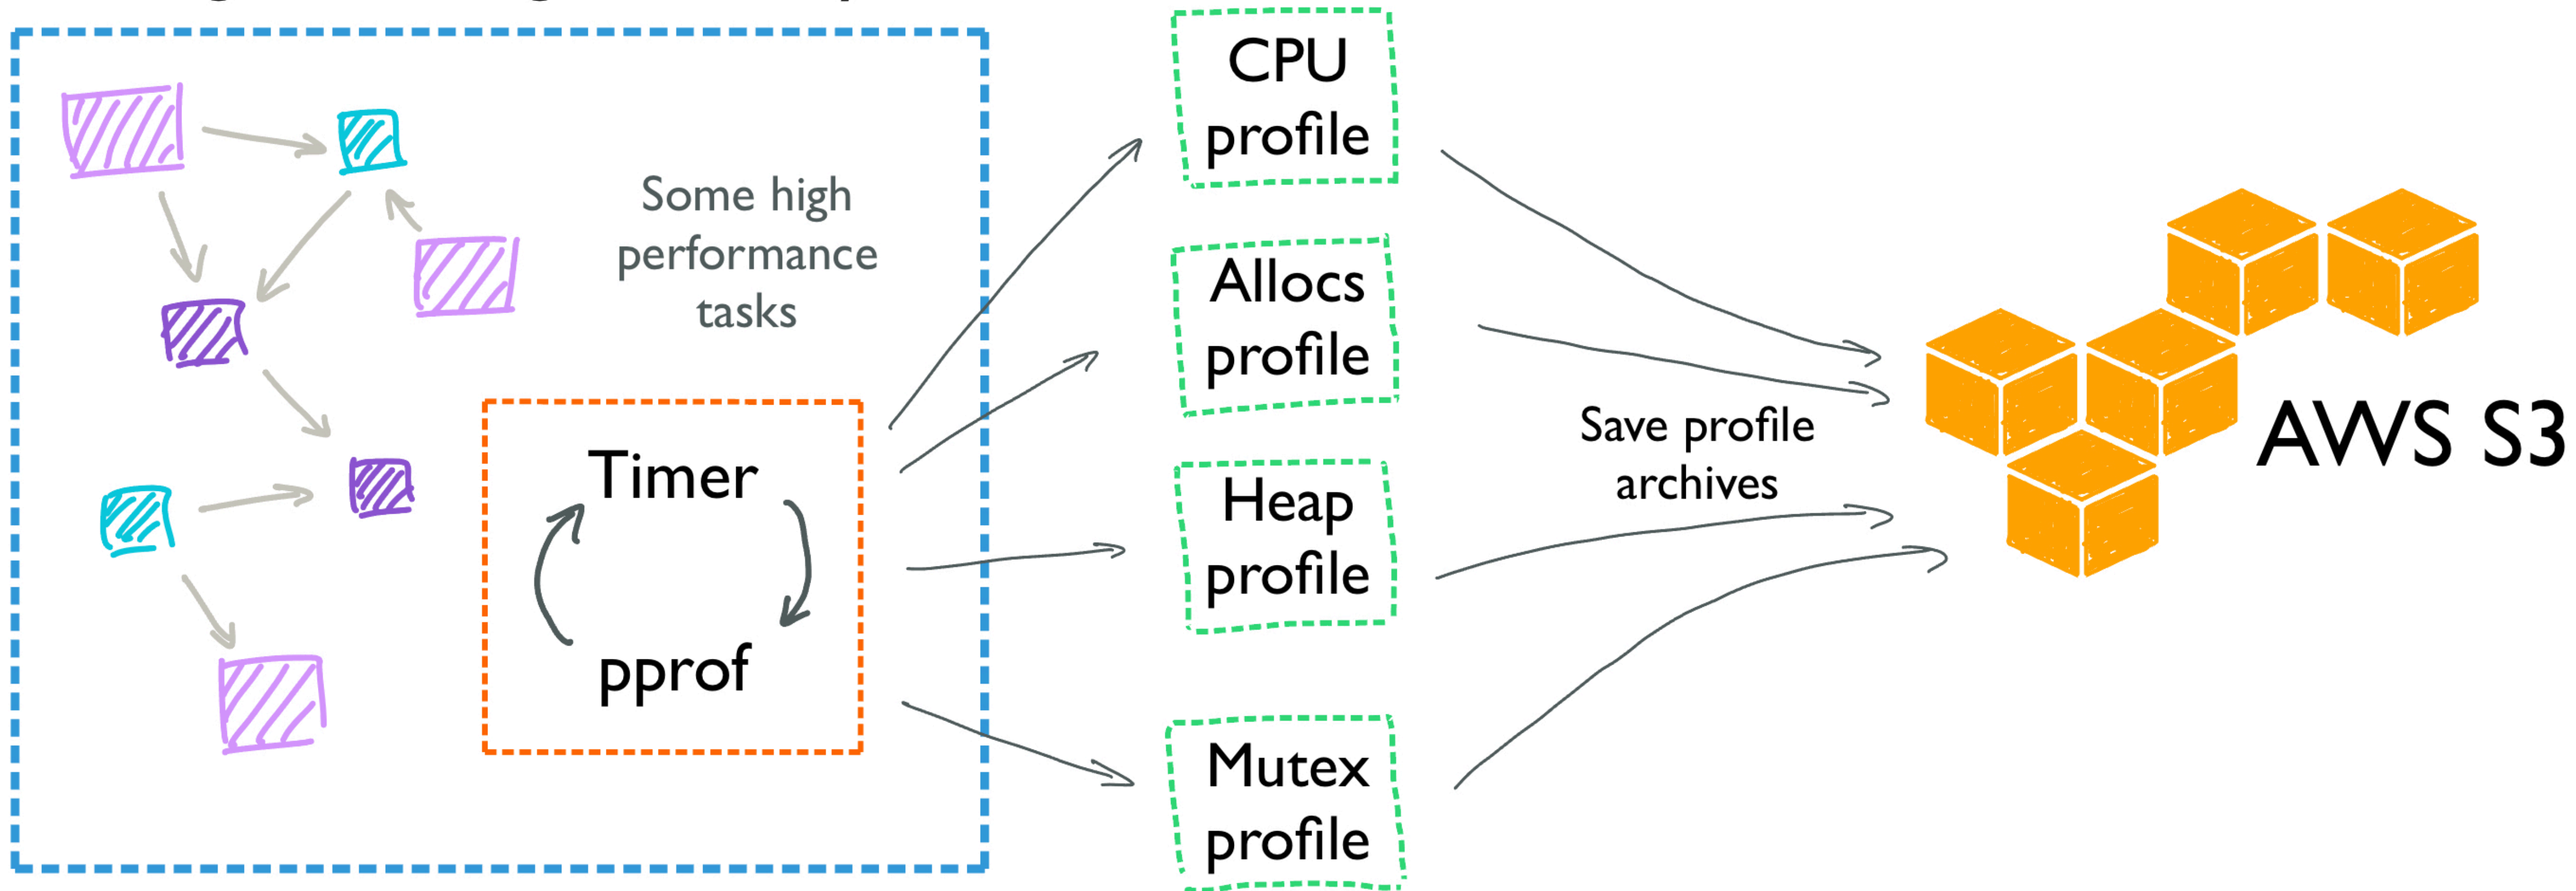
Possible ways to run pprof on runtime

- ▶ Run pprof HTTP server, and call specific pprof endpoints
 - ▶ The customer need to **expose a port** to internet on production machine
- ▶ Use **google/gops** to analyze agent program in runtime
 - ▶ The customer need to connect to container, take profile dumps and **share with us**
- ▶ Dump profiles through **go test**
 - ▶ Not applicable in the runtime

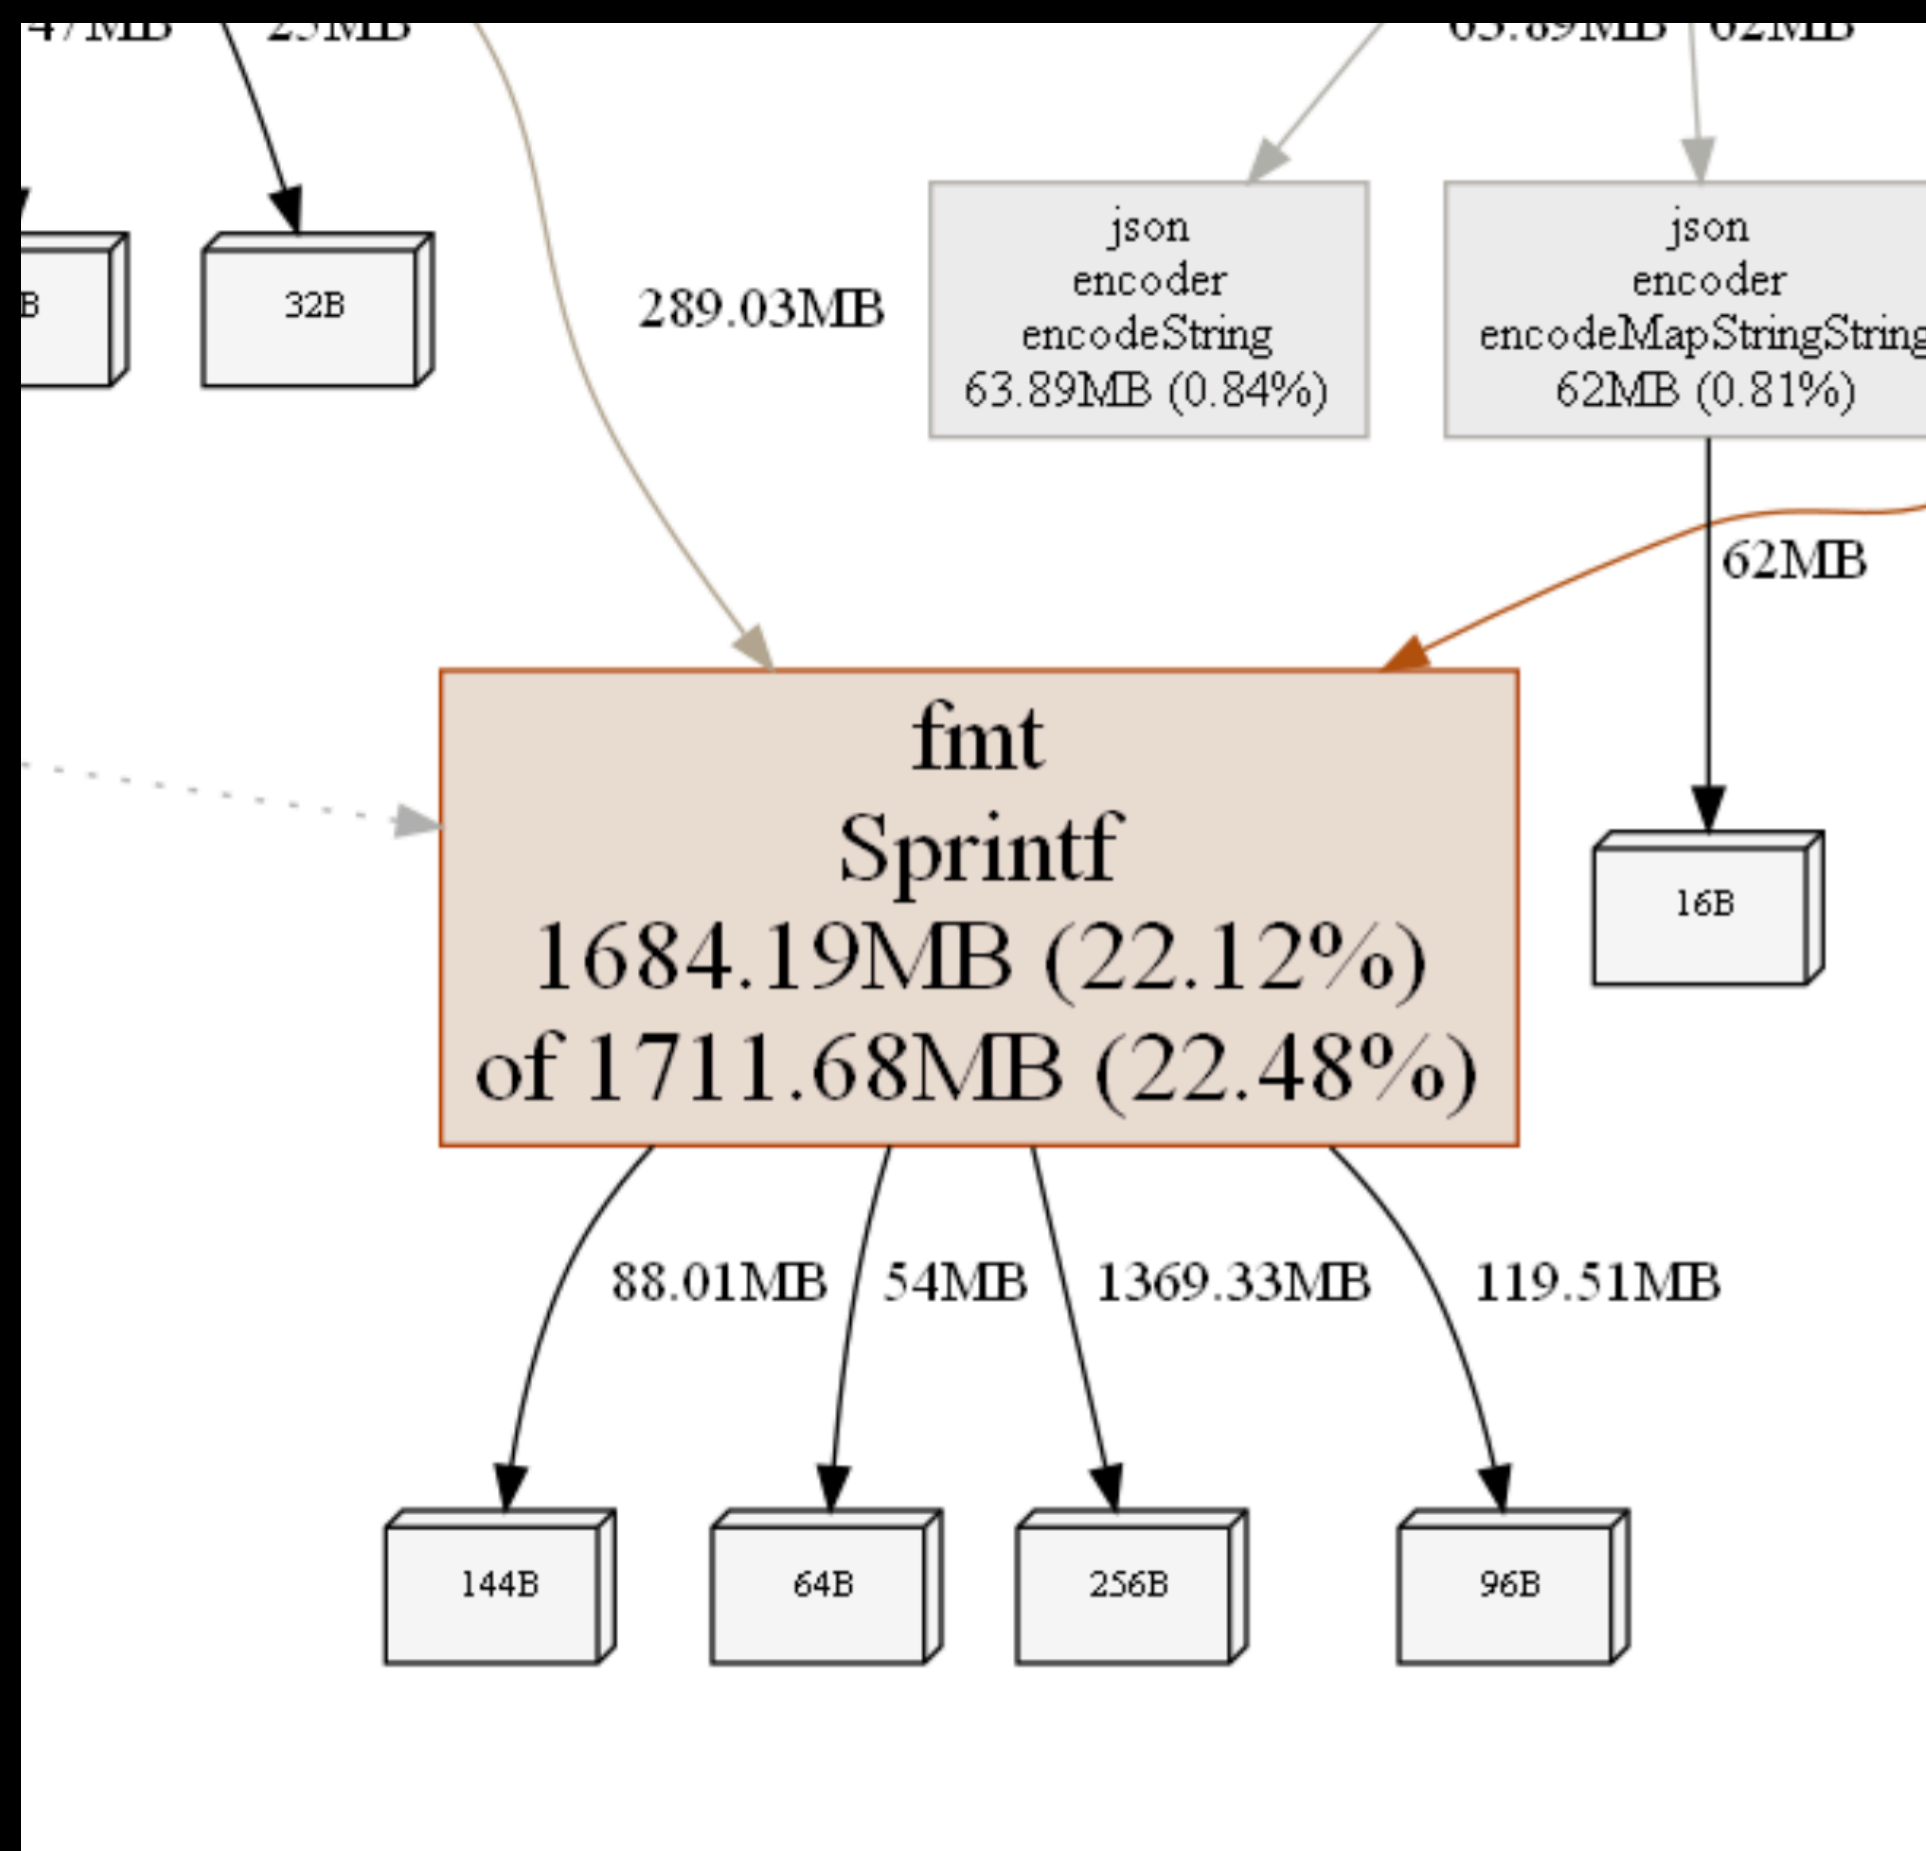
Solution: runtime/pprof

- ▶ Supports self-profiling on the agent binary runtime
- ▶ The pprof HTTP server actually uses runtime/pprof in background
- ▶ Supports various types of profiles
 - heap, allocs, goroutine, threadcreate, block, mutex, cpu

Edge Delta Agent Binary



It actually works!



We're hiring!



edgedelta.com/company/careers

sozan@edgedelta.com

Thank you for listening!